

iulsegboe

April 17, 2025

```
[118]: from snowflake.snowpark import Session
import pandas as pd
from snowflake.connector.pandas_tools import write_pandas
```

```
[75]: connection_parameters = {
    "account": "XOB39151",
    "user": "CHIPMUNK",
    "password": "m65LM2RqSfp696Y",
    "role": "TRAINING_ROLE",
    "warehouse": "ANIMAL_TASK_WH",
    "database": "NETFLIX_PROJECT_DB",
    "schema": "NETFLIX_ANALYTICS"
}
```

```
[77]: session = Session.builder.configs(connection_parameters).create()
```

```
[11]: global_alltime = session.table("CLEAN_GLOBAL_ALLTIME").to_pandas()
global_alltime.head()
```

```
[11]:
```

	SHOW_TITLE	CATEGORY	RANK	HOURS_VIEWED_FIRST_91_DAYS	\
0	red notice	films (english)	1	454200000	
1	don't look up	films (english)	2	408600000	
2	carry-on	films (english)	3	340800000	
3	the adam project	films (english)	4	281000000	
4	bird box	films (english)	5	325300000	

	VIEWS_FIRST_91_DAYS	RUNTIME	SEASON_TITLE
0	230900000	2	N/A
1	171400000	2	N/A
2	170400000	2	N/A
3	157600000	2	N/A
4	157400000	2	N/A

```
[15]: # Add a format column
global_alltime['FORMAT'] = global_alltime['CATEGORY'].str.contains('series',
↪case=False).astype(int)
```

```
[21]: global_alltime.columns = global_alltime.columns.str.lower()
```

```
[23]: global_alltime.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 40 entries, 0 to 39
Data columns (total 8 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   show_title            40 non-null    object
 1   category              40 non-null    object
 2   rank                  40 non-null    int8
 3   hours_viewed_first_91_days  40 non-null    int64
 4   views_first_91_days    40 non-null    int32
 5   runtime               40 non-null    int8
 6   season_title          40 non-null    object
 7   format                40 non-null    int64
dtypes: int32(1), int64(2), int8(2), object(3)
memory usage: 1.9+ KB
```

```
[25]: # Drop rows with nulls in key features
global_alltime = global_alltime.dropna(subset=['runtime',
↪ 'views_first_91_days'])

# Ensure numeric types
global_alltime['runtime'] = global_alltime['runtime'].astype(float)
global_alltime['views_first_91_days'] = global_alltime['views_first_91_days'].
↪ astype(float)
```

```
[27]: from sklearn.preprocessing import StandardScaler


features = global_alltime[['runtime', 'views_first_91_days', 'format']]
scaler = StandardScaler()
X_scaled = scaler.fit_transform(features)
```

```
[37]: inertias = []
for k in range(2, 10):
    km = KMeans(n_clusters=k, random_state=42)
    km.fit(X_scaled)
    inertias.append(km.inertia_)

# Plot Elbow Curve
import matplotlib.pyplot as plt

plt.figure(figsize=(8,5))
plt.plot(range(2,10), inertias, marker='o')
plt.title("Elbow Method for Optimal Clusters")
```

```
plt.xlabel("Number of clusters")
plt.ylabel("Inertia")
plt.grid(True)
plt.show()
```



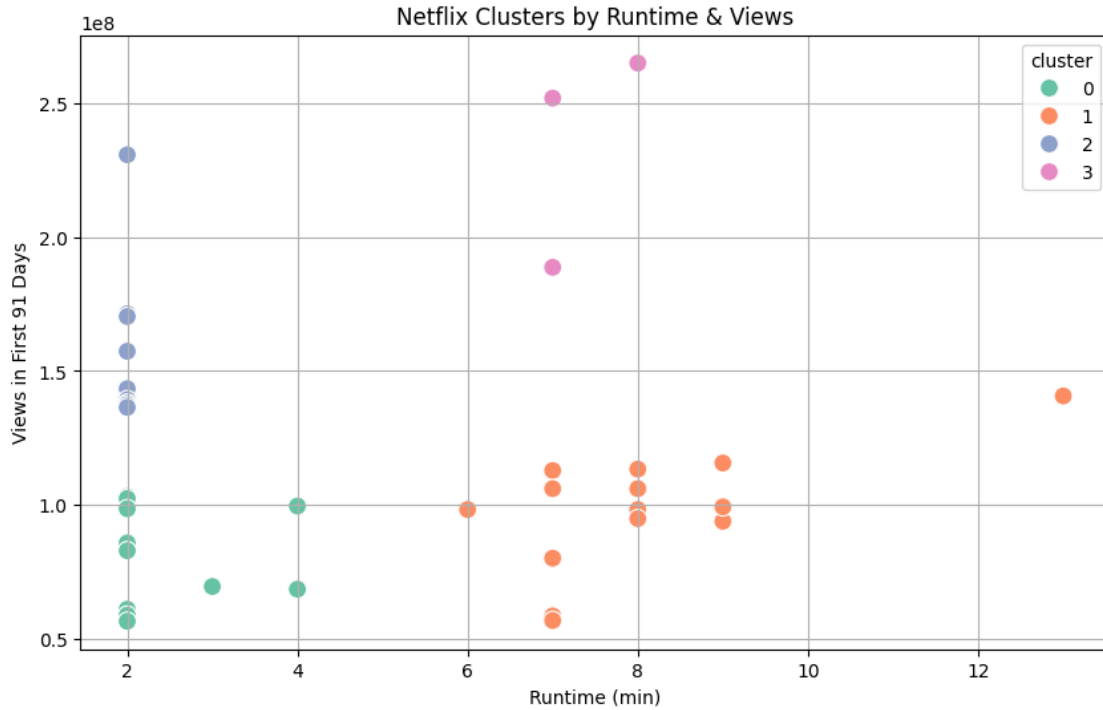
```
[29]: from sklearn.cluster import KMeans

kmeans = KMeans(n_clusters=4, random_state=42)
global_alltime['cluster'] = kmeans.fit_predict(X_scaled)
```

```
[31]: import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(10, 6))
sns.scatterplot(
    x=global_alltime['runtime'],
    y=global_alltime['views_first_91_days'],
    hue=global_alltime['cluster'],
    palette='Set2',
    s=100
)
plt.title("Netflix Clusters by Runtime & Views")
```

```
plt.xlabel("Runtime (min)")
plt.ylabel("Views in First 91 Days")
plt.grid(True)
plt.show()
```



```
[33]: from sklearn.metrics import silhouette_score

score = silhouette_score(X_scaled, global_alltime['cluster'])
print("Silhouette Score:", round(score, 3))
```

Silhouette Score: 0.599

```
[35]: session.write_pandas(
    global_alltime[['show_title', 'runtime', 'views_first_91_days', 'format', 'cluster']],
    "CLUSTERED_CONTENT",
    auto_create_table=True,
    overwrite=True
)
```

```
[35]: <snowflake.snowpark.table.Table at 0x13ff1a450>
```

```
[49]: global_alltime_regression = global_alltime.copy()
```

```

# Format
global_alltime_regression['format'] = global_alltime_regression['category'].str.
    ↪lower().str.contains('series').astype(int)

# Encode category
global_alltime_regression['category_encoded'] =
    ↪global_alltime_regression['category'].astype('category').cat.codes

# Convert runtime from hours to minutes
global_alltime['runtime'] = global_alltime['runtime'].astype(float) * 60

```

```

[51]: from sklearn.linear_model import LinearRegression
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import r2_score, mean_squared_error
      import numpy as np

      # Features and target
      X = global_alltime_regression[['runtime', 'format', 'category_encoded']]
      y = global_alltime_regression['views_first_91_days']

      # Split data
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
          ↪random_state=42)

      # Train model
      model = LinearRegression()
      model.fit(X_train, y_train)

      # Evaluate
      y_pred = model.predict(X_test)
      print("R² Score:", round(r2_score(y_test, y_pred), 3))
      rmse = np.sqrt(mean_squared_error(y_test, y_pred))
      print("RMSE:", round(rmse, 2))

```

R² Score: -0.602
 RMSE: 39958539.57

```

[61]: # Input values
      runtime = 100
      format = 1
      category_encoded = 3

      # Simulate
      new_data = pd.DataFrame([[runtime, format, category_encoded]], columns=X.
          ↪columns)
      predicted_views = model.predict(new_data)[0]

```

```
print("Predicted Views in 91 Days:", round(predicted_views))
```

Predicted Views in 91 Days: 687775229

```
[67]: simulated_inputs = pd.DataFrame([
      [40, 0, 1],
      [90, 1, 2],
      [120, 0, 3]
    ], columns=['runtime', 'format', 'category_encoded'])

    # Add predicted views
    simulated_inputs['predicted_views'] = model.predict(simulated_inputs)

    # Add scenario names
    simulated_inputs['scenario'] = ['Short Film', '90-min Series', 'Epic Action_
    ↪Film']

    # Reorder columns
    simulated = simulated_inputs[['scenario', 'runtime', 'format',
    ↪'category_encoded', 'predicted_views']]
```

```
[69]: session.write_pandas(
      simulated,
      "WHAT_IF_PREDICTIONS",
      auto_create_table=True,
      overwrite=True
    )
```

[69]: <snowflake.snowpark.table.Table at 0x168d3a3c0>

```
[79]: # Load table
    country_weekly = session.table("CLEAN_COUNTRY_WEEKLY").to_pandas()

    # lowercasing + fixing column names
    country_weekly.columns = country_weekly.columns.str.lower().str.replace(' ',
    ↪'_')

    country_weekly.head()
```

```
[79]:   country_name country_iso2   week category  weekly_rank \
0   Argentina          AR  2025-03-02   films           1
1   Argentina          AR  2025-03-02   films           2
2   Argentina          AR  2025-03-02   films           3
3   Argentina          AR  2025-03-02   films           4
4   Argentina          AR  2025-03-02   films           5

      show_title season_title  cumulative_weeks_in_top_10
```

0	counterattack	N/A	1
1	uncharted	N/A	2
2	a copenhagen love story	N/A	1
3	ticket to paradise	N/A	2
4	despicable me 3	N/A	7

```
[81]: # Make sure all column names are clean
country_weekly.columns = country_weekly.columns.str.lower().str.replace(' ', '_')

# Preview columns
print(country_weekly.columns)

# Optional: Rename to make it more standard
country_weekly.rename(columns={
    'country_name': 'country',
    'show_title': 'title',
    'week': 'date',
    'weekly_rank': 'rank'
}, inplace=True)

# Check for nulls or duplicates
print(country_weekly.isnull().sum())
print(f"Duplicates: {country_weekly.duplicated().sum()}")
```

```
Index(['country_name', 'country_iso2', 'week', 'category', 'weekly_rank',
      'show_title', 'season_title', 'cumulative_weeks_in_top_10'],
      dtype='object')
country              0
country_iso2        0
date                0
category            0
rank                0
title               0
season_title        0
cumulative_weeks_in_top_10  0
dtype: int64
Duplicates: 0
```

```
[83]: # Keep only Top 10 ranks
country_weekly = country_weekly[country_weekly['rank'] <= 10].copy()
```

```
[85]: country_weekly['date'] = pd.to_datetime(country_weekly['date'])
```

```
[91]: selected_countries = ['India', 'United States', 'Brazil', 'Morocco', 'United_
    ↪Arab Emirates']
```

```
country_filtered = country_weekly[country_weekly['country'].
↳isin(selected_countries)]
```

```
[93]: country_filtered = country_filtered.sort_values(by=['country', 'title', 'date'])
country_filtered.reset_index(drop=True, inplace=True)
```

```
[103]: country_filtered.tail()
```

```
[103]:
```

	country	country_iso2	date	category	rank	\
18200	United States	US	2023-10-22	tv	1	
18201	United States	US	2023-10-29	tv	3	
18202	United States	US	2023-11-05	tv	8	
18203	United States	US	2023-08-13	films	6	
18204	United States	US	2021-09-19	films	9	
...	
19195	United States	US	2025-03-02	tv	2	
19196	United States	US	2023-06-11	films	5	
19197	United States	US	2023-06-18	films	8	
19198	United States	US	2023-05-14	films	6	
19199	United States	US	2023-05-21	films	7	

	title	\
18200	the fall of the house of usher	
18201	the fall of the house of usher	
18202	the fall of the house of usher	
18203	the fast and the furious	
18204	the father who moves mountains	
...	...	
19195	zero day	
19196	zookeeper	
19197	zookeeper	
19198	¡que viva méxico!	
19199	¡que viva méxico!	

	season_title	\
18200	The Fall of the House of Usher: Limited Series	
18201	The Fall of the House of Usher: Limited Series	
18202	The Fall of the House of Usher: Limited Series	
18203	N/A	
18204	N/A	
...	...	
19195	Zero Day: Limited Series	
19196	N/A	
19197	N/A	
19198	N/A	
19199	N/A	

	cumulative_weeks_in_top_10
18200	2
18201	3
18202	4
18203	1
18204	1
...	...
19195	2
19196	1
19197	2
19198	1
19199	2

[1000 rows x 8 columns]

```
[105]: from prophet import Prophet
import pandas as pd
from tqdm import tqdm
```

```
[ ]: forecast_results = []

# Group the dataset
grouped = country_weekly.groupby(['country', 'title'])

for (country, title), group in tqdm(grouped, total=len(grouped)):
    try:
        # Must have at least 4 records for Prophet to work
        if group.shape[0] < 4:
            continue

        # Prepare time series format
        ts = group[['date', 'rank']].rename(columns={'date': 'ds', 'rank': 'y'})
        ts = ts.sort_values('ds')
        ts['y'] = 11 - ts['y'] # Reverse rank: higher = better

        # Fit the model
        model = Prophet()
        model.fit(ts)

        # Forecast 1 week ahead
        future = model.make_future_dataframe(periods=1, freq='W')
        forecast = model.predict(future)

        # Get last predicted value
        latest = forecast.tail(1)

        forecast_results.append({
```

```

        'country': country,
        'title': title,
        'forecast_week': latest['ds'].values[0],
        'predicted_score': latest['yhat'].values[0],
        'confidence_low': latest['yhat_lower'].values[0],
        'confidence_high': latest['yhat_upper'].values[0]
    })

except Exception as e:
    print(f"Error for {country} - {title}: {e}")
    continue

```

```

0%|                                     | 1/123788 [00:00<19:16:31,
1.78it/s]18:13:26 - cmdstanpy - INFO - Chain [1] start processing
18:13:27 - cmdstanpy - INFO - Chain [1] done processing
0%|                                     | 6/123788 [00:01<5:49:33,
5.90it/s]18:13:27 - cmdstanpy - INFO - Chain [1] start processing
18:13:27 - cmdstanpy - INFO - Chain [1] done processing
18:13:27 - cmdstanpy - INFO - Chain [1] start processing
18:13:27 - cmdstanpy - INFO - Chain [1] done processing
0%|                                     | 22/123788 [00:01<1:21:12,
25.40it/s]18:13:27 - cmdstanpy - INFO - Chain [1] start processing
18:13:27 - cmdstanpy - INFO - Chain [1] done processing
0%|                                     | 30/123788 [00:01<1:11:05,
29.01it/s]18:13:27 - cmdstanpy - INFO - Chain [1] start processing
18:13:27 - cmdstanpy - INFO - Chain [1] done processing
18:13:27 - cmdstanpy - INFO - Chain [1] start processing
18:13:27 - cmdstanpy - INFO - Chain [1] done processing
18:13:27 - cmdstanpy - INFO - Chain [1] start processing
18:13:27 - cmdstanpy - INFO - Chain [1] done processing
0%|                                     | 40/123788 [00:01<55:30,
37.16it/s]18:13:27 - cmdstanpy - INFO - Chain [1] start processing
18:13:27 - cmdstanpy - INFO - Chain [1] done processing
18:13:27 - cmdstanpy - INFO - Chain [1] start processing
18:13:27 - cmdstanpy - INFO - Chain [1] done processing
0%|                                     | 61/123788 [00:01<31:11,
66.11it/s]18:13:27 - cmdstanpy - INFO - Chain [1] start processing
18:13:27 - cmdstanpy - INFO - Chain [1] done processing
18:13:27 - cmdstanpy - INFO - Chain [1] start processing
18:13:27 - cmdstanpy - INFO - Chain [1] done processing
18:13:27 - cmdstanpy - INFO - Chain [1] start processing
18:13:27 - cmdstanpy - INFO - Chain [1] done processing
18:13:27 - cmdstanpy - INFO - Chain [1] start processing
18:13:27 - cmdstanpy - INFO - Chain [1] done processing
18:13:27 - cmdstanpy - INFO - Chain [1] start processing
18:13:27 - cmdstanpy - INFO - Chain [1] done processing
0%|                                     | 72/123788 [00:02<37:56,

```

```

18:54:10 - cmdstanpy - INFO - Chain [1] done processing
18:54:10 - cmdstanpy - INFO - Chain [1] start processing
18:54:10 - cmdstanpy - INFO - Chain [1] done processing
96%|          | 118551/123788 [40:44<01:01,
85.48it/s]18:54:10 - cmdstanpy - INFO - Chain [1] start processing
18:54:10 - cmdstanpy - INFO - Chain [1] done processing
18:54:10 - cmdstanpy - INFO - Chain [1] start processing
18:54:10 - cmdstanpy - INFO - Chain [1] done processing
18:54:10 - cmdstanpy - INFO - Chain [1] start processing
18:54:10 - cmdstanpy - INFO - Chain [1] done processing

```

```

[112]: forecast_df = pd.DataFrame(forecast_results)
forecast_df['predicted_rank'] = 11 - forecast_df['predicted_score']
forecast_df = forecast_df.sort_values(['country', 'predicted_rank'])

```

```

[114]: forecast_df.head()

```

```

[114]:      country      title forecast_week \
182  Argentina  the equalizer 2    2025-01-19
189  Argentina  the lincoln lawyer  2024-11-03
60   Argentina      emily in paris  2024-10-06
93   Argentina  john wick: chapter 3 - parabellum  2024-01-28
156  Argentina      sex education  2023-10-29

      predicted_score  confidence_low  confidence_high  predicted_rank
182          28.534371         28.508582         28.559863         -17.534371
189          26.791828         26.784899         26.798756         -15.791828
60           17.378035         15.997608         18.736663          -6.378035
93           11.326164         11.326162         11.326165          -0.326164
156           10.977782          7.824085         14.257803           0.022218

```

```

[122]: forecast_df['forecast_week'] = forecast_df['forecast_week'].dt.
      ↪strftime('%Y-%m-%d')

```

```

[124]: session.write_pandas(
    forecast_df,
    table_name="PREDICTED_TOP_SHOWS_NEXT_WEEK",
    database="NETFLIX_PROJECT_DB",
    schema="NETFLIX_ANALYTICS",
    auto_create_table=True,
    overwrite=True
)

```

```

/opt/anaconda3/lib/python3.12/site-packages/snowflake/snowpark/session.py:3132:
UserWarning: Pandas Dataframe has non-standard index of type <class
'pandas.core.indexes.base.Index'> which will not be written. Consider changing
the index to pd.RangeIndex(start=0,...,step=1) or call reset_index() to keep
index as column(s)

```

```
success, _, _, ci_output = write_pandas(
```

```
[124]: <snowflake.snowpark.table.Table at 0x166ca4980>
```

```
[128]: # Load CLEAN_COUNTRY_WEEKLY from Snowflake
country_weekly2 = session.table("CLEAN_COUNTRY_WEEKLY").to_pandas()
```

```
[130]: country_weekly2.head()
```

```
[130]:
```

	COUNTRY_NAME	COUNTRY_ISO2	WEEK	CATEGORY	WEEKLY_RANK	\
0	Argentina	AR	2025-03-02	films	1	
1	Argentina	AR	2025-03-02	films	2	
2	Argentina	AR	2025-03-02	films	3	
3	Argentina	AR	2025-03-02	films	4	
4	Argentina	AR	2025-03-02	films	5	

	SHOW_TITLE	SEASON_TITLE	CUMULATIVE_WEEKS_IN_TOP_10
0	counterattack	N/A	1
1	uncharted	N/A	2
2	a copenhagen love story	N/A	1
3	ticket to paradise	N/A	2
4	despicable me 3	N/A	7

```
[132]: # Rename and clean columns
country_weekly2.columns = country_weekly2.columns.str.lower().str.replace(" ",
↳ "_")
country_weekly2 = country_weekly2.sort_values(by=['country_name', 'show_title',
↳ 'week'])
country_weekly2.rename(columns={
    'country_name': 'country',
    'show_title': 'title',
    'week': 'date',
    'weekly_rank': 'rank'
}, inplace=True)

# Convert to datetime
country_weekly2['date'] = pd.to_datetime(country_weekly2['date'])

# SHIFT: Get next week's rank to create target
country_weekly2['next_rank'] = country_weekly2.groupby(['country',
↳ 'title'])['rank'].shift(-1)
country_weekly2['is_rank1_next_week'] = (country_weekly2['next_rank'] == 1).
↳ astype(int)

# Encode categorical fields
country_weekly2['country_encoded'] = country_weekly2['country'].
↳ astype('category').cat.codes
```

```

country_weekly2['title_encoded'] = country_weekly2['title'].astype('category').
    ↪cat.codes

# Feature: Trend (change in rank)
country_weekly2['prev_rank'] = country_weekly2.groupby(['country',
    ↪'title'])['rank'].shift(1)
country_weekly2['rank_trend'] = country_weekly2['prev_rank'] -
    ↪country_weekly2['rank']

# Feature: Weeks in Top 10
country_weekly2['weeks_in_top10'] = country_weekly2.groupby(['country',
    ↪'title']).cumcount() + 1

# Drop rows with NaNs (from shifting)
country_weekly2.dropna(subset=['next_rank', 'prev_rank'], inplace=True)

```

```

[136]: from xgboost import XGBClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report

features = ['rank', 'rank_trend', 'weeks_in_top10', 'country_encoded',
    ↪'title_encoded']
X = country_weekly2[features]
y = country_weekly2['is_rank1_next_week']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

model = XGBClassifier()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
print(classification_report(y_test, y_pred))

# Add predictions and probabilities
country_weekly2['rank1_prob'] = model.predict_proba(X)[: , 1]

```

	precision	recall	f1-score	support
0	0.95	0.99	0.97	28779
1	0.71	0.23	0.35	2075
accuracy			0.94	30854
macro avg	0.83	0.61	0.66	30854
weighted avg	0.93	0.94	0.93	30854

```
[138]: latest_date = country_weekly2['date'].max()
latest = country_weekly2[country_weekly2['date'] == latest_date].copy()

# Top predicted show per country
top_preds = latest.sort_values(['country', 'rank1_prob'], ascending=[True,
↪False]) \
                .groupby('country').first().reset_index()

# Clean result
top_preds = top_preds[['country', 'title', 'date', 'rank1_prob']]
top_preds.rename(columns={'date': 'forecast_week'}, inplace=True)

# Optional: round probability
top_preds['rank1_prob'] = top_preds['rank1_prob'].round(3)
```

```
[140]: from snowflake.connector.pandas_tools import write_pandas

# Get low-level connection from session
conn = session._conn._conn

write_pandas(
    conn=conn,
    df=top_preds,
    table_name="PREDICTED_RANK1_NEXT_WEEK",
    database="NETFLIX_PROJECT_DB",
    schema="NETFLIX_ANALYTICS",
    auto_create_table=True,
    overwrite=True
)
```

```
[140]: (True,
1,
16,
[('kookxvqkd1/file0.txt', 'LOADED', 16, 16, 1, 0, None, None, None, None)])
```

```
[ ]:
```